

Project 3: Local Feature Matching

CS 4476/6476: Computer Vision

Overview

The goal of this assignment is to create a local feature matching algorithm using techniques described in Szeliski chapter 4.1. The pipeline we suggest is a simplified version of the famous **SIFT** pipeline. The matching pipeline is intended to work for *instance-level* matching -- multiple views of the same physical scene.

Setup

1. Install **Miniconda**. It doesn't matter whether you use 2.7 or 3.6 because we will create our own environment anyways.
2. Create a conda environment, using the appropriate command. On Windows, open the installed "Conda prompt" to run this command. On MacOS and Linux, you can just use a terminal window to run the command. Modify the command based on your OS ('linux', 'mac', or 'win'):

```
conda env create -f proj3_env_<OS>.yaml
```
3. This should create an environment named 'proj3'. Activate it using the following Windows command:

```
activate proj3
```

or the following MacOS / Linux command:

```
source activate proj3
```
4. Install the project package, by running `pip install -e .` inside the repo folder.
5. Run the notebook using:

```
jupyter notebook ./proj3_code/proj3.ipynb
```
6. Generate the submission once you've finished the project using:

```
python zip_submission.py
```

Details

For this project, you need to implement the three major steps of a local feature matching algorithm:

- Interest point detection in `student_harris.py` (see Szeliski 4.1.1)
 - Local feature description in `student_sift.py` (see Szeliski 4.1.2)
 - Feature Matching in `student_feature_matching.py` (see Szeliski 4.1.3)
- There are numerous papers in the computer vision literature addressing each stage. For this project, we will suggest specific, relatively simple algorithms for each stage. You are encouraged to experiment with more sophisticated algorithms!

Interest point detection (`student_harris.py`)

You will implement the Harris corner detector as described in the lecture materials and Szeliski 4.1.1. See Algorithm 4.1 in the textbook for pseudocode. The starter code gives some additional suggestions. The main function will be `get_interest_points()` while implementing `my_filter2D()`, `get_gradients()`, `get_gaussian_kernel()`, `second_moments()`, `corner_response()`, `non_max_suppression()` will be helper

functions that will have tests to check progress as you work through Harris corner detection. You do not need to worry about scale invariance or keypoint orientation estimation for your baseline Harris corner detector. The original paper by Chris Harris and Mike Stephens describing their corner detector can be found [here](#).

Local feature description (`student_sift.py`)

You will implement a SIFT-like local feature as described in the lecture materials and Szeliski 4.1.2. See `get_features()` for more details. If you want to get your matching pipeline working quickly (and maybe to help debug the other algorithm stages), you might want to start with normalized patches as your local feature. There are 2 helper functions in `student_sift.py` that you will need to code to get your SIFT pipeline working. `get_magnitudes_and_orientations()` will return the magnitudes and orientations of the image gradients at every pixel. `get_feat_vec()` will get the feature vector associated with a specific interest point. Once these are done, move on to coding `get_features()`, which will combine these to get feature vectors for all interest points. More info about each function can be found in the function headers.

Feature matching (`student_feature_matching.py`)

You will implement the "ratio test" or "nearest neighbor distance ratio test" method of matching local features as described in the lecture materials and Szeliski 4.1.3. See equation 4.18 in particular. The potential matches that pass the ratio test the easiest should have a greater tendency to be correct matches -- think about *why*.

Using the starter code (`proj3.ipynb`)

The top-level `proj3.ipynb` IPython notebook provided in the starter code includes file handling, visualization, and evaluation functions. The correspondence will be visualized with `show_correspondence_circles()` and `show_correspondence_lines()` (you can comment one or both out if you prefer).

For the Notre Dame image pair there is a ground truth evaluation in the starter code as well. `evaluate_correspondence()` will classify each match as correct or incorrect based on hand-provided matches (see `show_ground_truth_corr()` for details). The starter code also contains ground truth correspondences for two other image pairs (Mount Rushmore and Episcopal Gaudi). You can test on those images by uncommenting the appropriate lines at the top of `proj3.ipynb`.

As you implement your feature matching pipeline, you should see your performance according to `evaluate_correspondence()` increase. Hopefully you find this useful, but don't *overfit* to the initial Notre Dame image pair which is relatively easy. The baseline algorithm suggested here and in the starter code will give you full credit and work fairly well on these Notre Dame images, but additional image pairs provided in the folder `data` are more difficult. They might exhibit more viewpoint, scale, and illumination variation. If you add enough Bells & Whistles you should be able to match more difficult image pairs.

Suggested implementation strategy

It is **highly suggested** that you implement the functions in this order:

- First, use `cheat_interest_points()` instead of `get_interest_points()`. This function will only work for the 3 image pairs with ground truth correspondence. This function cannot be used in your final implementation. It directly loads interest points from the the ground truth correspondences for the test cases. Even with this cheating, your accuracy will initially be near zero because the starter code features are empty and the starter code matches don't exist. `cheat_interest_points()` might return non-integer values, but you'll have to cut patches out at integer coordinates. You should address this by truncating the numbers.
- Second, change `get_features()` to return a simple feature. Start with, for instance, 16x16 patches centered on each interest point. Image patches aren't a great feature (they're not invariant to brightness change, contrast change, or small spatial shifts) but this is simple to implement and provides a baseline. You won't see your accuracy increase yet because the placeholder code in `match_features()` isn't assigning matches.
- Third, implement `match_features()`. Accuracy should increase on the Notre Dame pair if you're using 16x16 (256 dimensional) patches as your feature and if you only evaluate your 100 most confident matches. Accuracy on the other test cases will be lower
- Fourth, finish `get_features()` by implementing a sift-like feature. Accuracy should increase to 70% on the Notre Dame pair, 40% on Mount Rushmore, and 15% on Episcopal Gaudi if you only evaluate your 100 most confident matches (these are just estimates). These accuracies still aren't great because the human selected keypoints from `cheat_interest_points()` might not match particularly well according to your feature.
- Fifth, stop using `cheat_interest_points()` and implement `get_interest_points()`. Harris corners aren't as good as ground-truth points which we know correspond, so accuracy may drop. On the other hand, you can get hundreds or even a few thousand interest points so you have more opportunities to find confident matches. If you only evaluate the most confident 100 matches (see the `num_pts_to_evaluate` parameter) on the Notre Dame pair, you should be able to achieve >80% accuracy. You will likely need to do extra credit to get high accuracy on Mount Rushmore and Episcopal Gaudi.

Potentially useful NumPy (Python library), OpenCV, and SciPy

functions: `np.arctan2()`, `np.sort()`, `np.reshape()`, `np.newaxis`, `np.argsort()`, `np.gradient()`, `np.histogram()`, `np.hypot()`, `np.fliplr()`, `np.flipud()`, `cv2.getGaussianKernel()`

Forbidden functions (you can use for testing, but not in your final code): `cv2.SIFT()`, `cv2.SURF()`, `cv2.BFMatcher()`, `cv2.BFMatcher().match()`, `cv2.FlannBasedMatcher().knnMatch()`, `cv2.BFMatcher().knnMatch()`, `cv2.HOGDescriptor()`, `cv2.cornerHarris()`, `cv2.FastFeatureDetector()`, `cv2.ORB()`, `skimage.feature`, `skimage.feature.hog()`, `skimage.feature.daisy`, `skimage.feature.corner_harris()`, `skimage.feature.corner_shi_tomasi()`, `skimage.feature.match_descriptors()`, `skimage.feature.ORB()`, `scipy.signal.convolve()`, `cv2.filter2D()`, `cv2.Sobel()`, .

We haven't enumerated all possible forbidden functions here but using anyone else's code

that performs filtering, interest point detection, feature computation, or feature matching for you is forbidden.

Unit Tests

To test your code, go the main directory and run the command `pytest`.

Tips, Tricks, and Common Problems

- Make sure you're not swapping x and y coordinates at some point. If your interest points aren't showing up where you expect or if you're getting out of bound errors you might be swapping x and y coordinates. Remember, images expressed as NumPy arrays are accessed `image[y, x]`.
- Make sure you're features aren't somehow degenerate. You can visualize your features with `plt.imshow(image1_features)`, although you may need to normalize them first. If the features are mostly zero or mostly identical you may have made a mistake.

Writeup

For the writeup, the specific visualizations and questions to answer are in the powerpoint template.

Bells and Whistles (Partially required for grad students.)

Implementation of bells and whistles can increase your grade by up to 20 points (potentially over 100). The max score for all students is 120. Grad students will be graded out of 110, so to get full credit they will need to attempt some extra credit.

For all extra credit, be sure to include quantitative analysis showing the impact of the particular method you've implemented. Each item is "up to" some amount of points because trivial implementations may not be worthy of full extra credit

- up to 5 pts: (On Mount Rushmore image) Try varying some of the hand selected variables and see how the accuracy of the feature matching changes. Specifically vary the standard deviation of the second moment filter with these options [3, 6, 10, 30] and report the images of the ground truth correspondence and their corresponding accuracies in the report. On top of this vary the feature width in SIFT by these options [8, 16, 24, 32] and report the ground truth correspondence images and the corresponding accuracies for each different value. For the experiment regarding SIFT feature width, you may need to adjust the window size in the function `remove_border_vals` to ensure all SIFT patches will remain entirely in the image. When changing these variables ensure all others are set to default.
 - **Report Question:** When changing the values for large sigma (>20), why are the accuracies generally the same?
 - **Report Question:** What is the significance of changing feature width in SIFT?

Local feature matching bells and whistles:

An issue with the baseline matching algorithm is the computational expense of computing distance between all pairs of features. For a reasonable implementation of the base pipeline, this is likely to be the slowest part of the code. There are numerous schemes to try and approximate or accelerate feature matching:

- up to 8 pts: Create a lower dimensional descriptor that is still accurate enough. For example, if the descriptor is 32 dimensions instead of 128 then the distance computation should be about 4 times faster. PCA would be a good way to create a low dimensional descriptor. You would need to compute the PCA basis on a sample of your local descriptors from many images. However, calling `sklearn.decomposition.PCA()` is not enough to get credit (you need to implement PCA yourself), but you are allowed to use `numpy.linalg.svd()`. The function for pca can be found in `feature_matching.py`
 - PCA is an algorithm that is based around Singular Value Decomposition where you take the SVD of the covariance matrix of your data and reduce the number of usable eigen vectors. Then we project the data onto the remaining usable eigenvectors. Here is a link of the Georgia Tech Machine Learning that explains it in more depth:
<https://www.youtube.com/watch?v=kw9RonD69OU> (there are 3 parts)
- up to 7 pts: Use a space partitioning data structure like a kd-tree or some third party approximate nearest neighbor package to accelerate matching. You must achieve a runtime of less than a second (just matching) for full credit (with 1500 interest points). Your accuracy must also remain > 80 % for Notre Dame. Code your implementation in `accelerated_matching()` in `student_feature_matching.py`.
 - **Report Question:** What did you try and what was the speedup? Why is it faster?

Rubric

- +30 pts: Implementation of Harris corner detector in `student_harris.py`
- +35 pts: Implementation of SIFT-like local feature in `student_sift.py`
- +10 pts: Implementation of "Ratio Test" matching in `student_feature_matching.py`
- +25 pts: Report.
- +20 pts: Extra credit (implementation of any bells & whistles up to twenty points, max score of 120).
- -5*n pts: Lose 5 points for every time you do not follow the instructions for the hand in format

Undergrad students will be graded out of 100, graduate students out of 110. Grad students will have to attempt some extra credit for full credit.

Handing in

- submission.zip created with `zip_submission`
- gt-username_proj3.pdf

Credits

Assignment developed by James Hays, Samarth Brahmabhatt, and John Lambert.
Updated by Judy Hoffman, Mitch Donley, and Vijay Upadhyia.